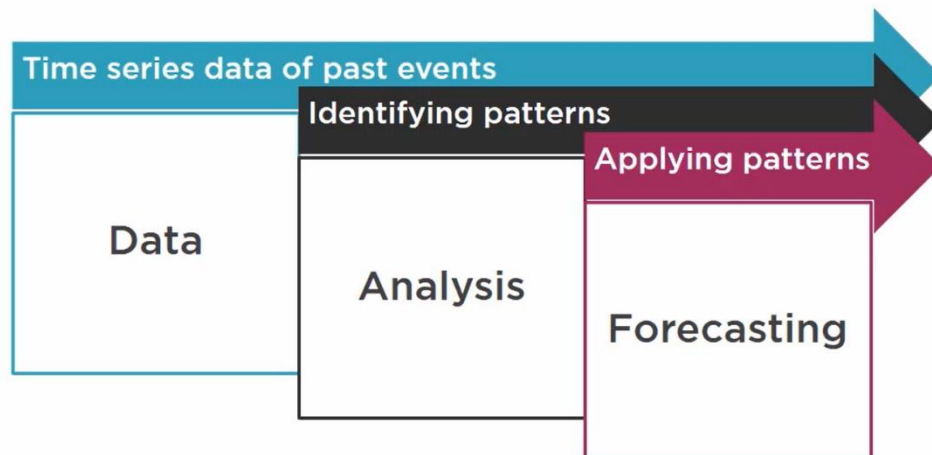


Time Series Analysis

Source: <https://app.pluralsight.com/library/courses/r-time-series-analysis-forecasting/table-of-contents>

Time Series Analysis Background:

How This Process Work





Univariate Time Series:

One variable attached to a time stamp.

1. Linear

- ARIMA
- Exponential Smoothing
- Simple Methods

2. Non-linear

- K Nearest Neighbors
- Clustering
- Neural Nets
- Support Vector Machines
- Q Learning
- Decision Trees

Multivariate Time Series:

Two or more variables attached to a time stamp.

Datasets:

1. Lynx trapping in Canada.
2. Temperature measurements in Nottingham.
3. Randomly generated series.

Time Series	Vector
The time stamp specifies a successive order for the values.	A unique ID does not necessarily provide a specific order to the data.

Converting Vectors to Time Series:

1. Functions `ts()`
 - Attaches a time stamp to a vector
 - Converts the class to 'ts'
 - Use it to build time series from scratch
2. Library (xts)
 - Importing time series data into R

Lag:

A gap between two or more observations.

Y_t : An observation of the time series

$Y_{114} = 3396$ (The last observation of 'lynx')

Lag of 1 = $Y_t - Y_{t-1} = Y_{114} - Y_{113} = 3396 - 2657$

Lag of 2 = $Y_t - Y_{t-2} = Y_{114} - Y_{112} = 3396 - 1590$


Lynx Dataset:

- mean()
- median()
- plot()
- sort()
- quantile()

Results:


```
> head(lynx)
[1] 269 321 585 871 1475 2821
> time(lynx)
Time Series:
Start = 1821
End = 1934
Frequency = 1
[1] 1821 1822 1823 1824 1825 1826 1827 1828 1829 1830 1831 1832 1833 1834 1835 1836 1837 1838 1839 1840 1841 1842 1843
[24] 1844 1845 1846 1847 1848 1849 1850 1851 1852 1853 1854 1855 1856 1857 1858 1859 1860 1861 1862 1863 1864 1865 1866
[47] 1867 1868 1869 1870 1871 1872 1873 1874 1875 1876 1877 1878 1879 1880 1881 1882 1883 1884 1885 1886 1887 1888 1889
[70] 1890 1891 1892 1893 1894 1895 1896 1897 1898 1899 1900 1901 1902 1903 1904 1905 1906 1907 1908 1909 1910 1911 1912
[93] 1913 1914 1915 1916 1917 1918 1919 1920 1921 1922 1923 1924 1925 1926 1927 1928 1929 1930 1931 1932 1933 1934
> length(lynx)
[1] 114
> mean(lynx); median(lynx)
[1] 1538.018
[1] 771
> plot(lynx)
> sort(lynx)
[1] 39 45 49 59 68 73 80 81 98 105 108 151 153 184 188 201 213 225
[19] 229 229 236 245 255 269 279 299 299 321 345 358 360 361 377 377 382 387
[37] 389 399 409 409 469 473 485 523 529 546 552 585 587 662 674 684 687 731
[55] 736 756 758 784 808 871 957 1000 1033 1132 1292 1307 1388 1426 1475 1537 1590 1594
[73] 1623 1638 1676 1824 1836 2042 2119 2129 2251 2285 2432 2511 2536 2577 2657 2685 2713 2725
[91] 2811 2821 2871 2935 2985 3091 3311 3396 3409 3465 3495 3574 3790 3794 3800 3928 4031 4254
[109] 4431 4950 5943 6313 6721 6991
> quantile(lynx)
 0%    25%    50%    75%   100%
39.00 348.25 771.00 2566.75 6991.00
> quantile(lynx, prob = seq(0, 1, length = 11), type = 5)
 0%   10%   20%   30%   40%   50%   60%   70%   80%   90%  100%
39.0 146.7 259.2 380.5 546.6 771.0 1470.1 2165.6 2818.0 3790.4 6991.0
>
```

Plot:



ts and mts:


```
Console Terminal x Background Jobs x
R 4.2.1 · ~/
> #random uniform data between 10 to 45
> mydata = runif(n = 50, min = 10, max = 45)
> #packing into a quarterly time series
> mytimeseries = ts(data = mydata, start = c(1956,3), frequency = 4)
> plot(mytimeseries)
> class(mydata)
[1] "numeric"
> class(mytimeseries)
[1] "ts"
> class(lynx)
[1] "ts"
> #a typical mts data set
> class(EuStockMarkets); head(EuStockMarkets)
[1] "mts"      "ts"      "matrix"
      DAX    SMI    CAC    FTSE
[1,] 1628.75 1678.1 1772.8 2443.6
[2,] 1613.63 1688.5 1750.5 2460.2
[3,] 1606.51 1678.6 1718.0 2448.2
[4,] 1621.04 1684.1 1708.1 2470.4
[5,] 1618.16 1686.6 1723.1 2484.7
[6,] 1610.61 1671.6 1714.3 2466.8
> |
```



Pattern to Identify:

- | | |
|------------------|---|
| 1. Trend: | Dataset moving towards a direction. |
| 2. Seasonality: | Repeated pattern over a fixed interval. |
| 3. Mean: | The average of the dataset. |
| 4. Variance: | Indicator of variability. |
| 5. Stationarity: | Constant mean and variance. |

Understanding Pattern Using Graph:




Seasonality

Constant variance

Clear trend

Non-stationary

Autocorrelation



Trend with exponential curve


Changing mean

Changing variance


Non-stationary

Transformation is required

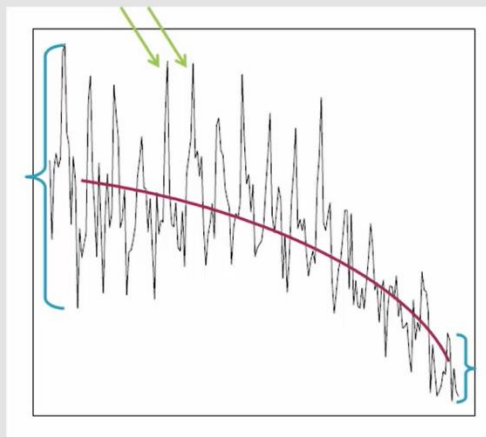
Exponential seasonality
Exponential trend
Changing variance
Non-stationary



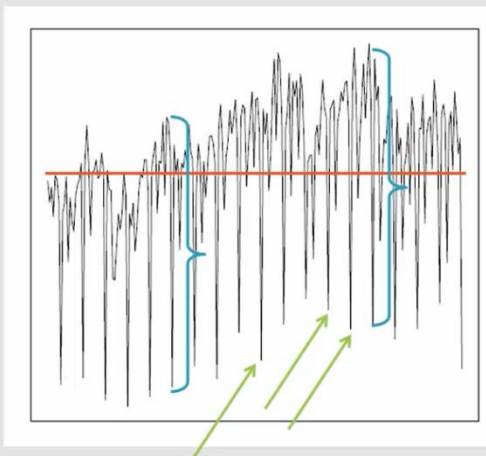
Constant mean
Changing variance
Heteroscedastic dataset
Non-stationary
Preprocessing prior analysis




Seasonality
Changing variance
Trend
Non-stationary



Seasonality
Constant mean
Constant variance
Non-stationary
Autocorrelation is present





Time series of random normally distributed data


No trend

Constant mean

Constant variance

Stationarity is present

Transformation and differencing is not needed



Time series with a clear trend


Increasing mean

Non-stationary dataset


Preprocessing prior modeling

Nottem Dataset:


- `plot()`



- `decompose()`




- `plot.ts()`



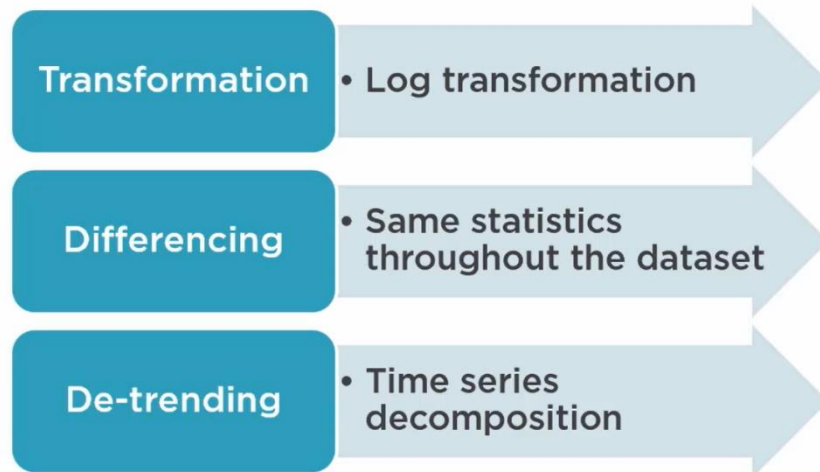
- `ggseasonplot()`

Seasonal plot: nottem



Stationarity:

What to Do with Non-stationary Data?



Console Terminal x Background Jobs x

R 4.2.1 · ~/

```
> #Stationary
> test = ts(c(rnorm(100,2,1), rnorm(100,50,1)),start = 1)
> plot(test)
> plot(diff(test))
> #Unit Root Tests
> x = rnorm(1000) #random normal data
> library(tseries)
> adf.test(x)
```

Augmented Dickey-Fuller Test

```
data: x
Dickey-Fuller = -10.397, Lag order = 9, p-value = 0.01
alternative hypothesis: stationary
```


```
Warning message:
In adf.test(x) : p-value smaller than printed p-value
> plot(nottem) #Seasonal data
> adf.test(nottem)
```

Augmented Dickey-Fuller Test


```
data: nottem
Dickey-Fuller = -12.998, Lag order = 6, p-value = 0.01
alternative hypothesis: stationary
```

```
Warning message:
In adf.test(nottem) : p-value smaller than printed p-value
> |
```


Test plot:



Diff(test) Plot:



Notttem Plot:




Auto-Correlation:


- `acf()` Shows the autocorrelation.
- `pacf()` Shows the partial autocorrelation.

Autocorrelation	Partial autocorrelation
The correlation coefficient between lags of the time series.	The correlation coefficient adjusted for shorter lags.

acf Vs pacf:


- Using lynx






- Using rnorm

Series rnorm(500)







Methods with Library Forecast:

Naïve Method	Seasonal naïve method	Mean method	Drift method
Returns the last observation as forecast value – <code>naive()</code>	Returns the last observation of the seasonal stage – <code>snaive()</code>	Returns the mean as forecast value – <code>meanf()</code>	Carries the change over first to last observation into the future – <code>rwf()</code>

```
Console Terminal Background Jobs
R 4.2.1 · ~/
> set.seed(50)
> testts <- ts(rnorm(300), start = c(1919,1), frequency = 4)
> plot(testts)
> library(forecast)
> meanmodel <- meanf(testts, h = 20)
> naivemodel <- naive(testts, h = 20)
> driftmodel <- rwf(testts, h = 20, drift = T)
> plot(meanmodel, main = "")
> lines(naivemodel$mean, col = 123, lwd = 2)
> lines(driftmodel$mean, col = 22, lwd = 2)
> legend("topleft", lty = 1, cex = 0.5, col = c(4, 123, 22), legend = c("Mean method", "Naive method", "Drift method"))
> |
```





Error Indicators:

- | | |
|-----------------------------------|------|
| • Mean Absolute Error- | MAE |
| • Root Mean Squared Error- | RMSE |
| • Mean Absolute Scale Error- | MASE |
| • Mean Absolute Percentage Error- | MAPE |

MAE:

$$\text{MAE} = \frac{\sum_{i=1}^n |e_i|}{n} = \frac{\sum_{i=1}^n |y_i - \hat{y}_i|}{n}$$

y_i	\hat{y}_i	$ y_i - \hat{y}_i = e_i $
2.45	2.63	$ -0.18 $
3.12	2.99	$ 0.13 $
3.06	3.15	$ -0.09 $
2.63	2.99	$ -0.36 $
1.89	2.21	$ -0.32 $
1.12	1.76	$ -0.64 $
2.56	1.43	$ 1.13 $

- ◀ Scale Dependent Errors: MAE
- ◀ The mean of all differences between actual and forecasted absolute values
- ◀ $|e_i|$ = absolute value of a forecast error
- ◀ $|e_i| = |y_i - \hat{y}_i|$
 y_i = actual value
 \hat{y}_i = forecast of y_i

$$\leftarrow \frac{\sum_{i=1}^n |e_i|}{n} = \frac{2.85}{7} = 0.407 = \text{MAE}$$

RMSE:

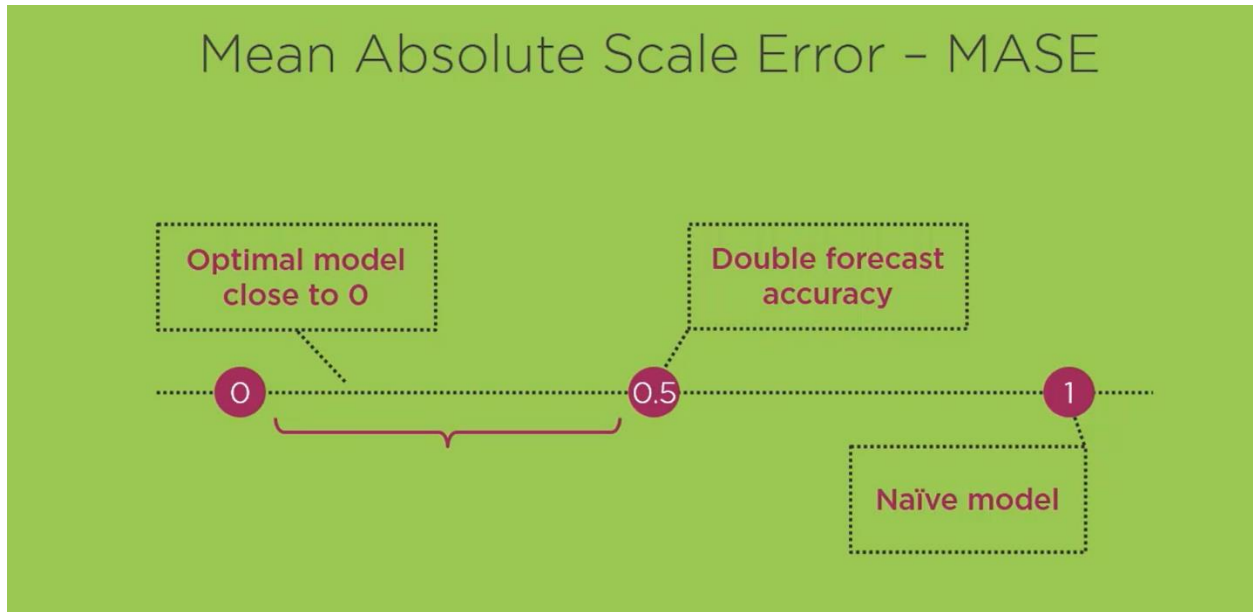
$$\text{RMSE} = \sqrt{\frac{\sum_{i=1}^n e_i^2}{n}} = \sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{n}}$$

y_i	\hat{y}_i	$(y_i - \hat{y}_i)^2 = e_i^2$
2.45	2.63	-0.18^2
3.12	2.99	0.13^2
3.06	3.15	0.09^2
2.63	2.99	-0.36^2
1.89	2.21	-0.32^2
1.12	1.76	-0.64^2
2.56	1.43	1.13^2

- ◀ Scale Dependent Errors: RMSE
- ◀ Standard deviation of differences between actual and forecasted values
- ◀ e_i^2 = squared value of a forecast error
- ◀ $e_i^2 = (y_i - \hat{y}_i)^2$
 y_i = actual value
 \hat{y}_i = forecast of y_i

$$\leftarrow \sqrt{\frac{\sum_{i=1}^n e_i^2}{n}} = \sqrt{\frac{1.976}{7}} = 0.531 = \text{RMSE}$$

MASE:





MAPE:

$$\text{MAPE} = \frac{\sum_{i=1}^n |p_i|}{n} = \frac{\sum_{i=1}^n \left| \frac{100e_i}{y_i} \right|}{n}$$

- ◀ Scale Independent Error: MAPE
- ◀ Measures the difference of forecast errors and divides it by the actual observation value
- ◀ $|p_i|$ = absolute value of forecast error differences
 y_i = actual value
- ◀ Does not allow for 0 values
- ◀ Puts more weight on extreme values and positive errors
- ◀ Use it to compare models on different datasets

Comparison:





```

Console Terminal x Background Jobs x
R 4.2.1 · ~/
> accuracy(meanmodel, mytstest)
              ME      RMSE      MAE      MPE      MAPE      MASE      ACF1 Theil's U
Training set  4.146651e-18  0.9922591  0.7999341  103.27287  103.27287  0.7247526  0.06882037    NA
Test set     -5.947488e-05  0.9250726  0.7350389   99.97963   99.97963  0.6659566 -0.15127726  0.9658077
> accuracy(naivemodel, mytstest)
              ME      RMSE      MAE      MPE      MAPE      MASE      ACF1 Theil's U
Training set -0.002185465  1.351649  1.086030 -32.96660  564.0232  0.9839598 -0.4830987    NA
Test set      1.714805941  1.948414  1.768427  97.41156  466.8504  1.6022227 -0.1512773  1.675834
> accuracy(driftmodel, mytstest)
              ME      RMSE      MAE      MPE      MAPE      MASE      ACF1 Theil's U
Training set -1.014813e-17  1.351647  1.086125 -32.43850  563.4943  0.9840464 -0.4830987    NA
Test set      1.805503e+00  2.032968  1.854351  95.38722  492.5072  1.6800705 -0.1268056  1.767978
>


```

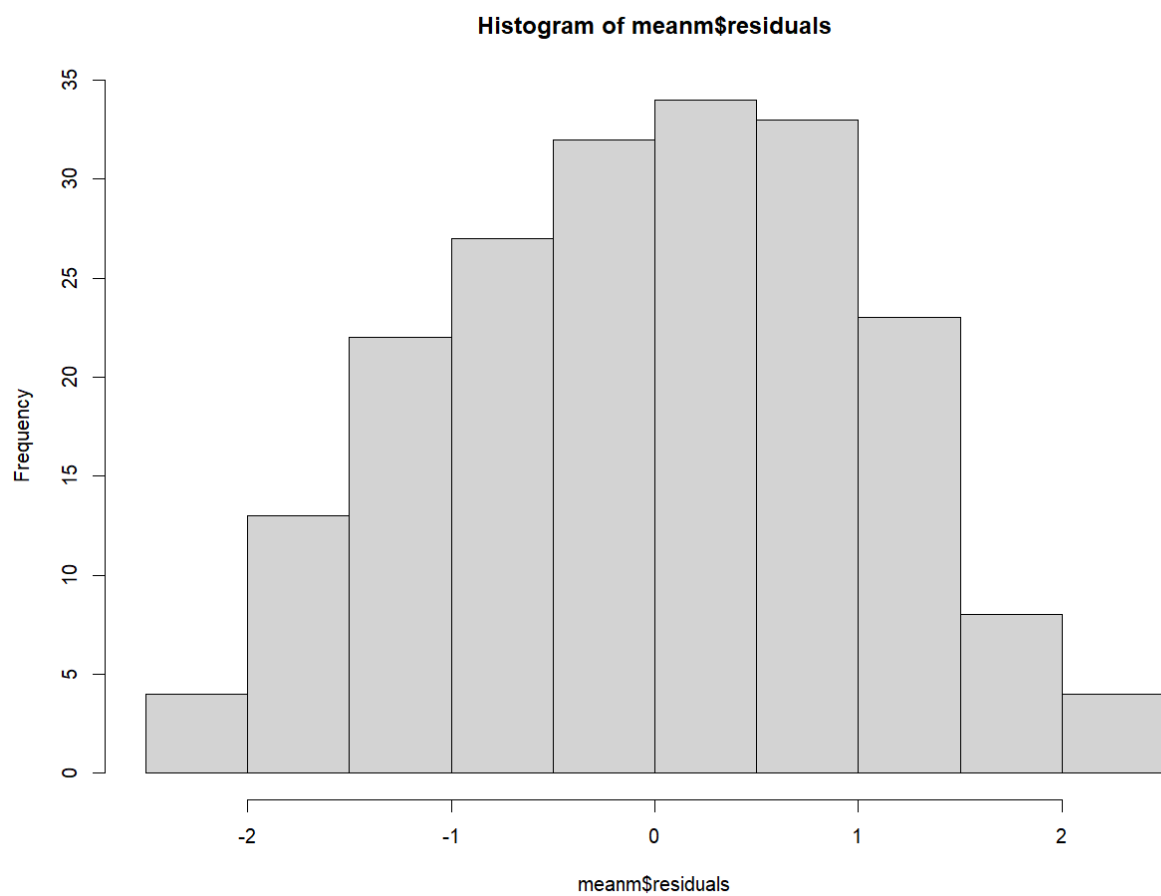
Residuals:


Ideal Model

Zero mean	Constant variance	Correlated residuals	Normal distribution
Fix: addition or subtraction	Fix: transformation – not always possible.	Fix: differencing	Fix: transformation – not always possible

```
Console Terminal x Background Jobs x
R 4.2.1 · ~/
> #Data set
> set.seed(95)
> myts <- ts(rnorm(200), start = (1919))
> #Setting up simple models
> library(forecast)
> meanm <- meanf(myts, h = 20)
> naivem <- naive(myts, h = 20)
> driftm <- rwf(myts, h = 20, drift = T)
> #Variance and mean of the mean model
> var(meanm$residuals)
[1] 1.053807
> plot(meanm$residuals)
> mean(meanm$residuals)
[1] -5.95498e-18
> #Deleting the NA at the front of the vector
> naiwithoutNA <- naivem$residuals
> naiwithoutNA <- naiwithoutNA[2:200]
> var(naiwithoutNA)
[1] 1.798592
> mean(naiwithoutNA)
[1] 0.006605028
> driftwithoutNA <- driftm$residuals
> driftwithoutNA <- driftwithoutNA[2:200]
> var(driftwithoutNA)
[1] 1.798592
> mean(driftwithoutNA)
[1] -4.502054e-17
> # Histogram of distribution
> hist(meanm$residuals)
> #Autocorrelation
> acf(meanm$residuals)
> |
```







The Mean Model on Random Data:

Zero mean	Normal distribution	Equal variance	No autocorrelation
mean()	hist()	var() and plot()	acf()

ARIMA:

ARIMA => Univariate => Non-seasonal

AR	<ul style="list-style-type: none"> Autoregressive term – “p”
I	<ul style="list-style-type: none"> Integration / differencing – “d”
MA	<ul style="list-style-type: none"> Moving average – “q”

Stationarity:

ARIMA(p, d, q)	ARMA(p, d)
Non-stationary time series gets differenced (“d”) before “p” and “q” get specified.	With stationary time series the autoregressive (“p”) and moving average (“q”) terms get ordered without differencing.

Variations of the Model:

AR(1) – ARIMA(1,0,0)	MA(1) – ARIMA (0,0,1)
Autoregressive model (“p” only)	Moving average model (“q” only)

What do the Parameters do?

p	Summation of lags – AR $Y_t = c + \varphi_1 * y_{t-1} + \varphi_2 * y_{t-2} + \dots + \varphi_p * y_{t-p}$
d	Degree of differencing – I
q	Summation of forecast error terms – MA $Y_t = c + \vartheta_1 * e_{t-1} + \vartheta_2 * e_{t-2} + \dots + \vartheta_q * e_{t-q}$

How to Calculate an AR Model

Coefficients:

	ar1	mean
	1.1246	1547.3859
s.e.	0.0903	136.8501

$$Y_t = c + \varphi_1 * y_{t-1}$$

Coefficients:

	ar1	ar2	ar3	ar4	mean
	1.1246	-0.7174	0.2634	-0.2543	1547.3859
s.e.	0.0903	0.1367	0.1361	0.0897	136.8501

$$Y_t = c + \varphi_1 * y_{t-1} + \varphi_2 * y_{t-2} + \varphi_3 * y_{t-3} + \varphi_4 * y_{t-4}$$

How to Calculate an ARMA Model

Coefficients:


	ar1	ma1	mean
	1.3421	-0.2027	1544.4039
s.e.	0.0984	0.1261	131.9242

$$Y_t = c + \varphi_1 * y_{t-1} + \vartheta_1 * e_{t-1}$$

ARIMA(0, 1, 0)


Drift: $c = Y_t - Y_{t-1}$

No drift: $Y_t = Y_{t-1}$




ARIMA Model Practice:

- p Autocorrelation is clear
- d It might be stationary
- q There might be forecasting errors



Series lynx



```

Console Terminal x Background Jobs x
R 4.2.1 · ~/
> plot(lynx)
> acf(lynx) ; pacf(lynx)
> auto.arima(lynx)
Series: lynx
ARIMA(2,0,2) with non-zero mean

Coefficients:
      ar1      ar2      ma1      ma2      mean
    1.3421 -0.6738 -0.2027 -0.2564 1544.4039
s.e.  0.0984  0.0801  0.1261  0.1097  131.9242

sigma^2 = 761965: log likelihood = -932.08
AIC=1876.17 AICc=1876.95 BIC=1892.58
> auto.arima(lynx, trace = T)

ARIMA(2,0,2) with non-zero mean : 1876.952
ARIMA(0,0,0) with non-zero mean : 2006.724
ARIMA(1,0,0) with non-zero mean : 1927.209
ARIMA(0,0,1) with non-zero mean : 1918.165
ARIMA(0,0,0) with zero mean : 2080.721
ARIMA(1,0,2) with non-zero mean : 1888.757
ARIMA(2,0,1) with non-zero mean : 1880.014
ARIMA(3,0,2) with non-zero mean : 1878.603
ARIMA(2,0,3) with non-zero mean : Inf
ARIMA(1,0,1) with non-zero mean : 1891.442
ARIMA(1,0,3) with non-zero mean : 1890.03
ARIMA(3,0,1) with non-zero mean : 1881.962
ARIMA(3,0,3) with non-zero mean : Inf
ARIMA(2,0,2) with zero mean : 1905.595

Best model: ARIMA(2,0,2) with non-zero mean

Series: lynx
ARIMA(2,0,2) with non-zero mean

Coefficients:
      ar1      ar2      ma1      ma2      mean
    1.3421 -0.6738 -0.2027 -0.2564 1544.4039
s.e.  0.0984  0.0801  0.1261  0.1097  131.9242

sigma^2 = 761965: log likelihood = -932.08
AIC=1876.17 AICc=1876.95 BIC=1892.58
> myar = auto.arima(lynx, stepwise = F, approximation = F)
> myar
Series: lynx
ARIMA(4,0,0) with non-zero mean

Coefficients:
      ar1      ar2      ar3      ar4      mean
    1.1246 -0.7174  0.2634 -0.2543 1547.3859
s.e.  0.0903  0.1367  0.1361  0.0897  136.8501

sigma^2 = 748457: log likelihood = -931.11
AIC=1874.22 AICc=1875.01 BIC=1890.64
> plot(forecast(myar, h = 3))
>

```

Calculating the ARIMA Model


```
Series: lynx
ARIMA(2,0,2) with non-zero mean

Coefficients:
      ar1      ar2      ma1      ma2      mean
1.3421 -0.6738 -0.2027 -0.2564 1544.4039
s.e.  0.0984  0.0801  0.1261  0.1097  131.9242

sigma^2 estimated as 761965:  log likelihood=-932.08
AIC=1876.17  AICc=1876.95  BIC=1892.58
```

$$Y_t = 1554.4 + 1.3421 * Y_{t-1} + (-0.6738) * Y_{t-2} + (-0.2027) * e_{t-1} + (-0.2564) * e_{t-2}$$

Forecasts from ARIMA(4,0,0) with non-zero mean

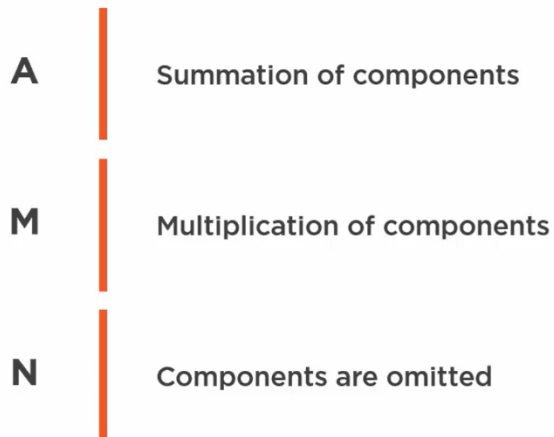


Exponential Smoothing

Parameters of Exponential Smoothing



Parameter Operators




Exponential Smoothing Function:

Function ses() Simple exponential smoothing	Function holt() Trend methods
Function hw() Holt-Winter seasonal method	Function ets() Selects the optimal model


ETS:

The 'model=' Argument of 'ets()'

- Z Auto Selection
- A Additive model
- M Multiplicative model
- N Non-present (except error)



Forecasts from ETS(A,N,A)



```

Console Terminal x Background Jobs x
R 4.2.1 · ~/
> # ets
> library(forecast)
> #using function ets
> etsmodel = ets(nottem) ; etsmodel
ETS(A,N,A)

Call:
ets(y = nottem)

Smoothing parameters:
  alpha = 0.0392
  gamma = 1e-04

Initial states:
  l = 49.4597
  s = -9.5635 -6.6186 0.5447 7.4811 11.5783 12.8567
      8.9762 3.4198 -2.7516 -6.8093 -9.7583 -9.3556

sigma: 2.3203

      AIC      AICc      BIC
1734.944 1737.087 1787.154
> #Plotting the model vs original
> plot(nottem, lwd = 3)
> lines(etsmodel$fitted, col = "red")
> #Plotting the forecast
> plot(forecast(etsmodel, h = 12))
> #Changing the prediction interval
> plot(forecast(etsmodel, h = 12, level = 95))
> #Manually setting the ets model
> etsmodmult = ets(nottem, model = "MZM")
> etsmodmult
ETS(M,N,M)

Call:
ets(y = nottem, model = "MZM")


Smoothing parameters:
  alpha = 0.0214
  gamma = 1e-04

Initial states:
  l = 49.3793
  s = 0.8089 0.8647 1.0132 1.1523 1.2348 1.2666
      1.1852 1.0684 0.9405 0.8561 0.8005 0.8088


sigma: 0.0508

      AIC      AICc      BIC
1761.911 1764.054 1814.121
> #Plot as comparison
> plot(nottem, lwd = 3)
> lines(etsmodmult$fitted, col = "red")
> |


```



Comparing the Models



Model 'ANA'



Model 'MNM'